

Natural Language Processing



Lecture 7: Parsing with Context Free Grammars II.
CKY for PCFGs. Earley Parser.

11/13/2020

COMS W4705
Yassine Benajiba

Recall: Syntactic Ambiguity

$S \rightarrow NP VP$	$NP \rightarrow she$
$VP \rightarrow V NP$	$NP \rightarrow glasses$
$VP \rightarrow VP PP$	$D \rightarrow the$
$PP \rightarrow P NP$	$N \rightarrow cat$
$NP \rightarrow D N$	$N \rightarrow glasses$
$NP \rightarrow NP PP$	$V \rightarrow saw$
	$P \rightarrow with$



Which parse tree is “better”? More probable?

Probabilities for Parse Trees

- Let \mathcal{T}_G be the set of all parse trees generated by grammar G .
- We want a model that assigns a probability to each parse tree, such that $\sum_{t \in \mathcal{T}_G} P(t) = 1$.
- We can use this model to select the most probable parse tree compatible with an input sentence.
 - This is another example of a generative model!

Selecting Parse Trees

- Let $\mathcal{T}_G(s)$ be the set of trees generated by grammar G whose *yield* (sequence of leafs) is string s .
- The most likely parse tree produced by G for string s is

$$\arg \max_{t \in \mathcal{T}_G(s)} P(t)$$

- How do we define $P(t)$?
- How do we learn such a model from training data (annotated or un-annotated).
- How do we find the highest probability tree for a given sentence? (*parsing/decoding*)

Probabilistic Context Free Grammars (PCFG)

- A PCFG consists of a Context Free Grammar $G=(N, \Sigma, R, S)$ and a probability $P(A \rightarrow \beta)$ for each production $A \rightarrow \beta \in R$.
- The probabilities for all rules with the same left-hand-side sum up to 1:

$$\sum_{A \rightarrow \beta: A=X} P(A \rightarrow \beta) = 1 \text{ for all } X \in N$$

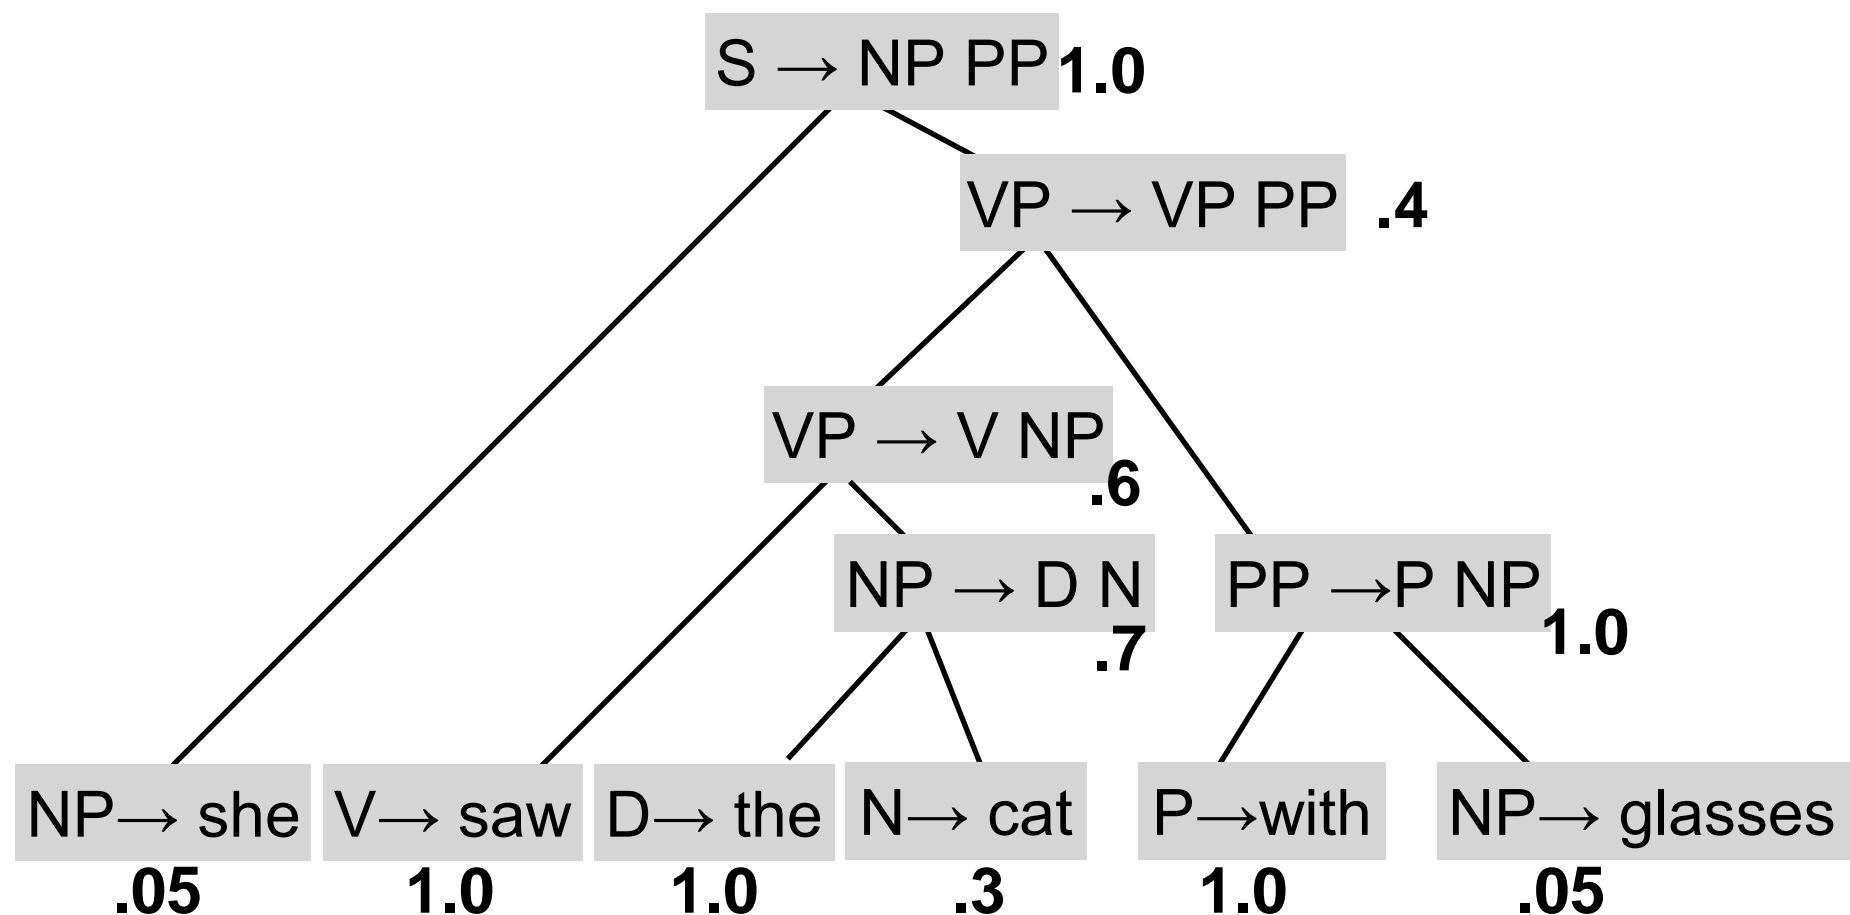
- Think of this as the conditional probability for $A \rightarrow \beta$, given the left-hand-side nonterminal A .

PCFG Example

S	→	NP VP	[1.0]	NP	→	she	[0.05]
VP	→	V NP	[0.6]	NP	→	glasses	[0.05]
VP	→	VP PP	[0.4]	D	→	the	[1.0]
PP	→	P NP	[1.0]	N	→	cat	[0.3]
NP	→	D N	[0.7]	N	→	glasses	[0.7]
NP	→	NP PP	[0.2]	V	→	saw	[1.0]
				P	→	with	[1.0]

Parse Tree Probability

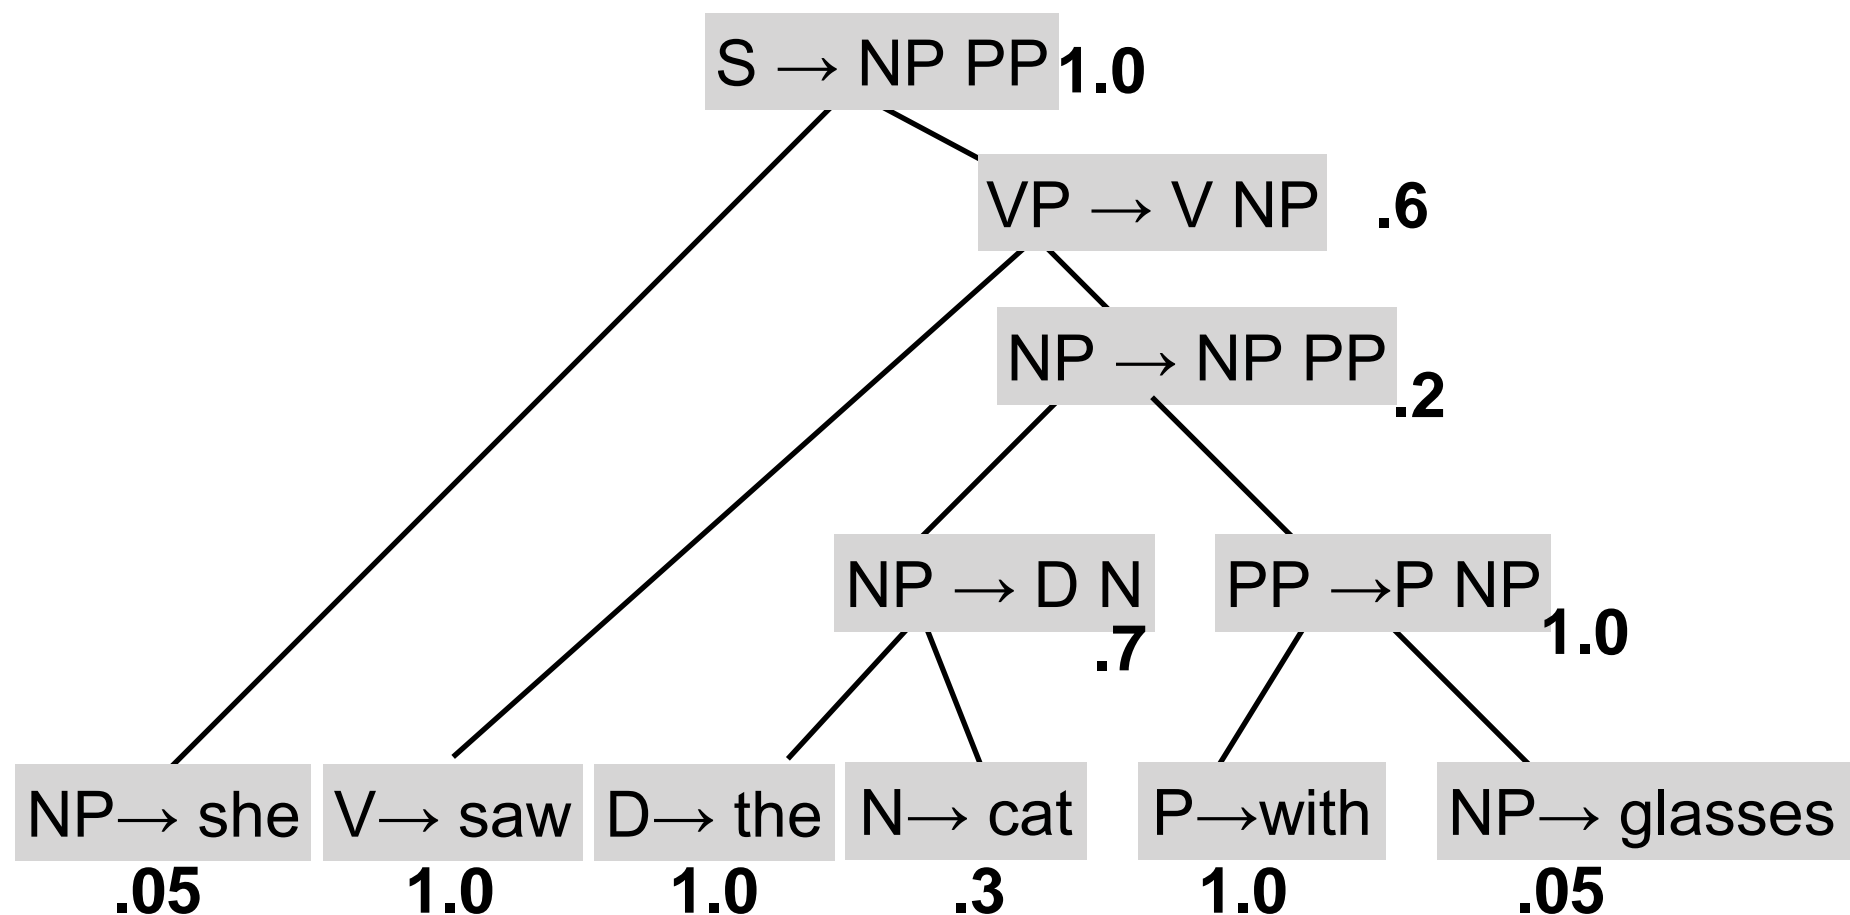
- Given a parse tree $t \in \mathcal{T}_G$, containing rules $A_1 \rightarrow \beta_1, \dots, A_n \rightarrow \beta_n$ the probability of t is
$$P(t) = \prod_{i=1}^n P(A_i \rightarrow \beta_i)$$



$$1 \times .05 \times .4 \times .6 \times 1 \times 0.7 \times 1 \times 0.3 \times 1 \times 1 \times .05 = .000126$$

Parse Tree Probability

- Given a parse tree $t \in \mathcal{T}_G$, containing rules $A_1 \rightarrow \beta_1, \dots, A_n \rightarrow \beta_n$ the probability of t is
$$P(t) = \prod_{i=1}^n P(A_i \rightarrow \beta_i)$$



$$1 \times .05 \times .6 \times 1 \times .2 \times .7 \times 1 \times .3 \times 1 \times 1 \times .05 = 0.000063 < 0.000126$$

Estimating PCFG probabilities

- Supervised training: We can estimate PCFG probabilities from a *treebank*, a corpus manually annotated with constituency structure using maximum likelihood estimates:

$$P(A \rightarrow \beta) = \frac{\text{count}(A \rightarrow \beta)}{\text{count}(A)}$$

- Unsupervised training:
 - What if we have a grammar and a corpus, but no annotated parses?
 - Can use the **inside-outside** algorithm for parsing and do EM estimation of the probabilities (not discussed in this course)

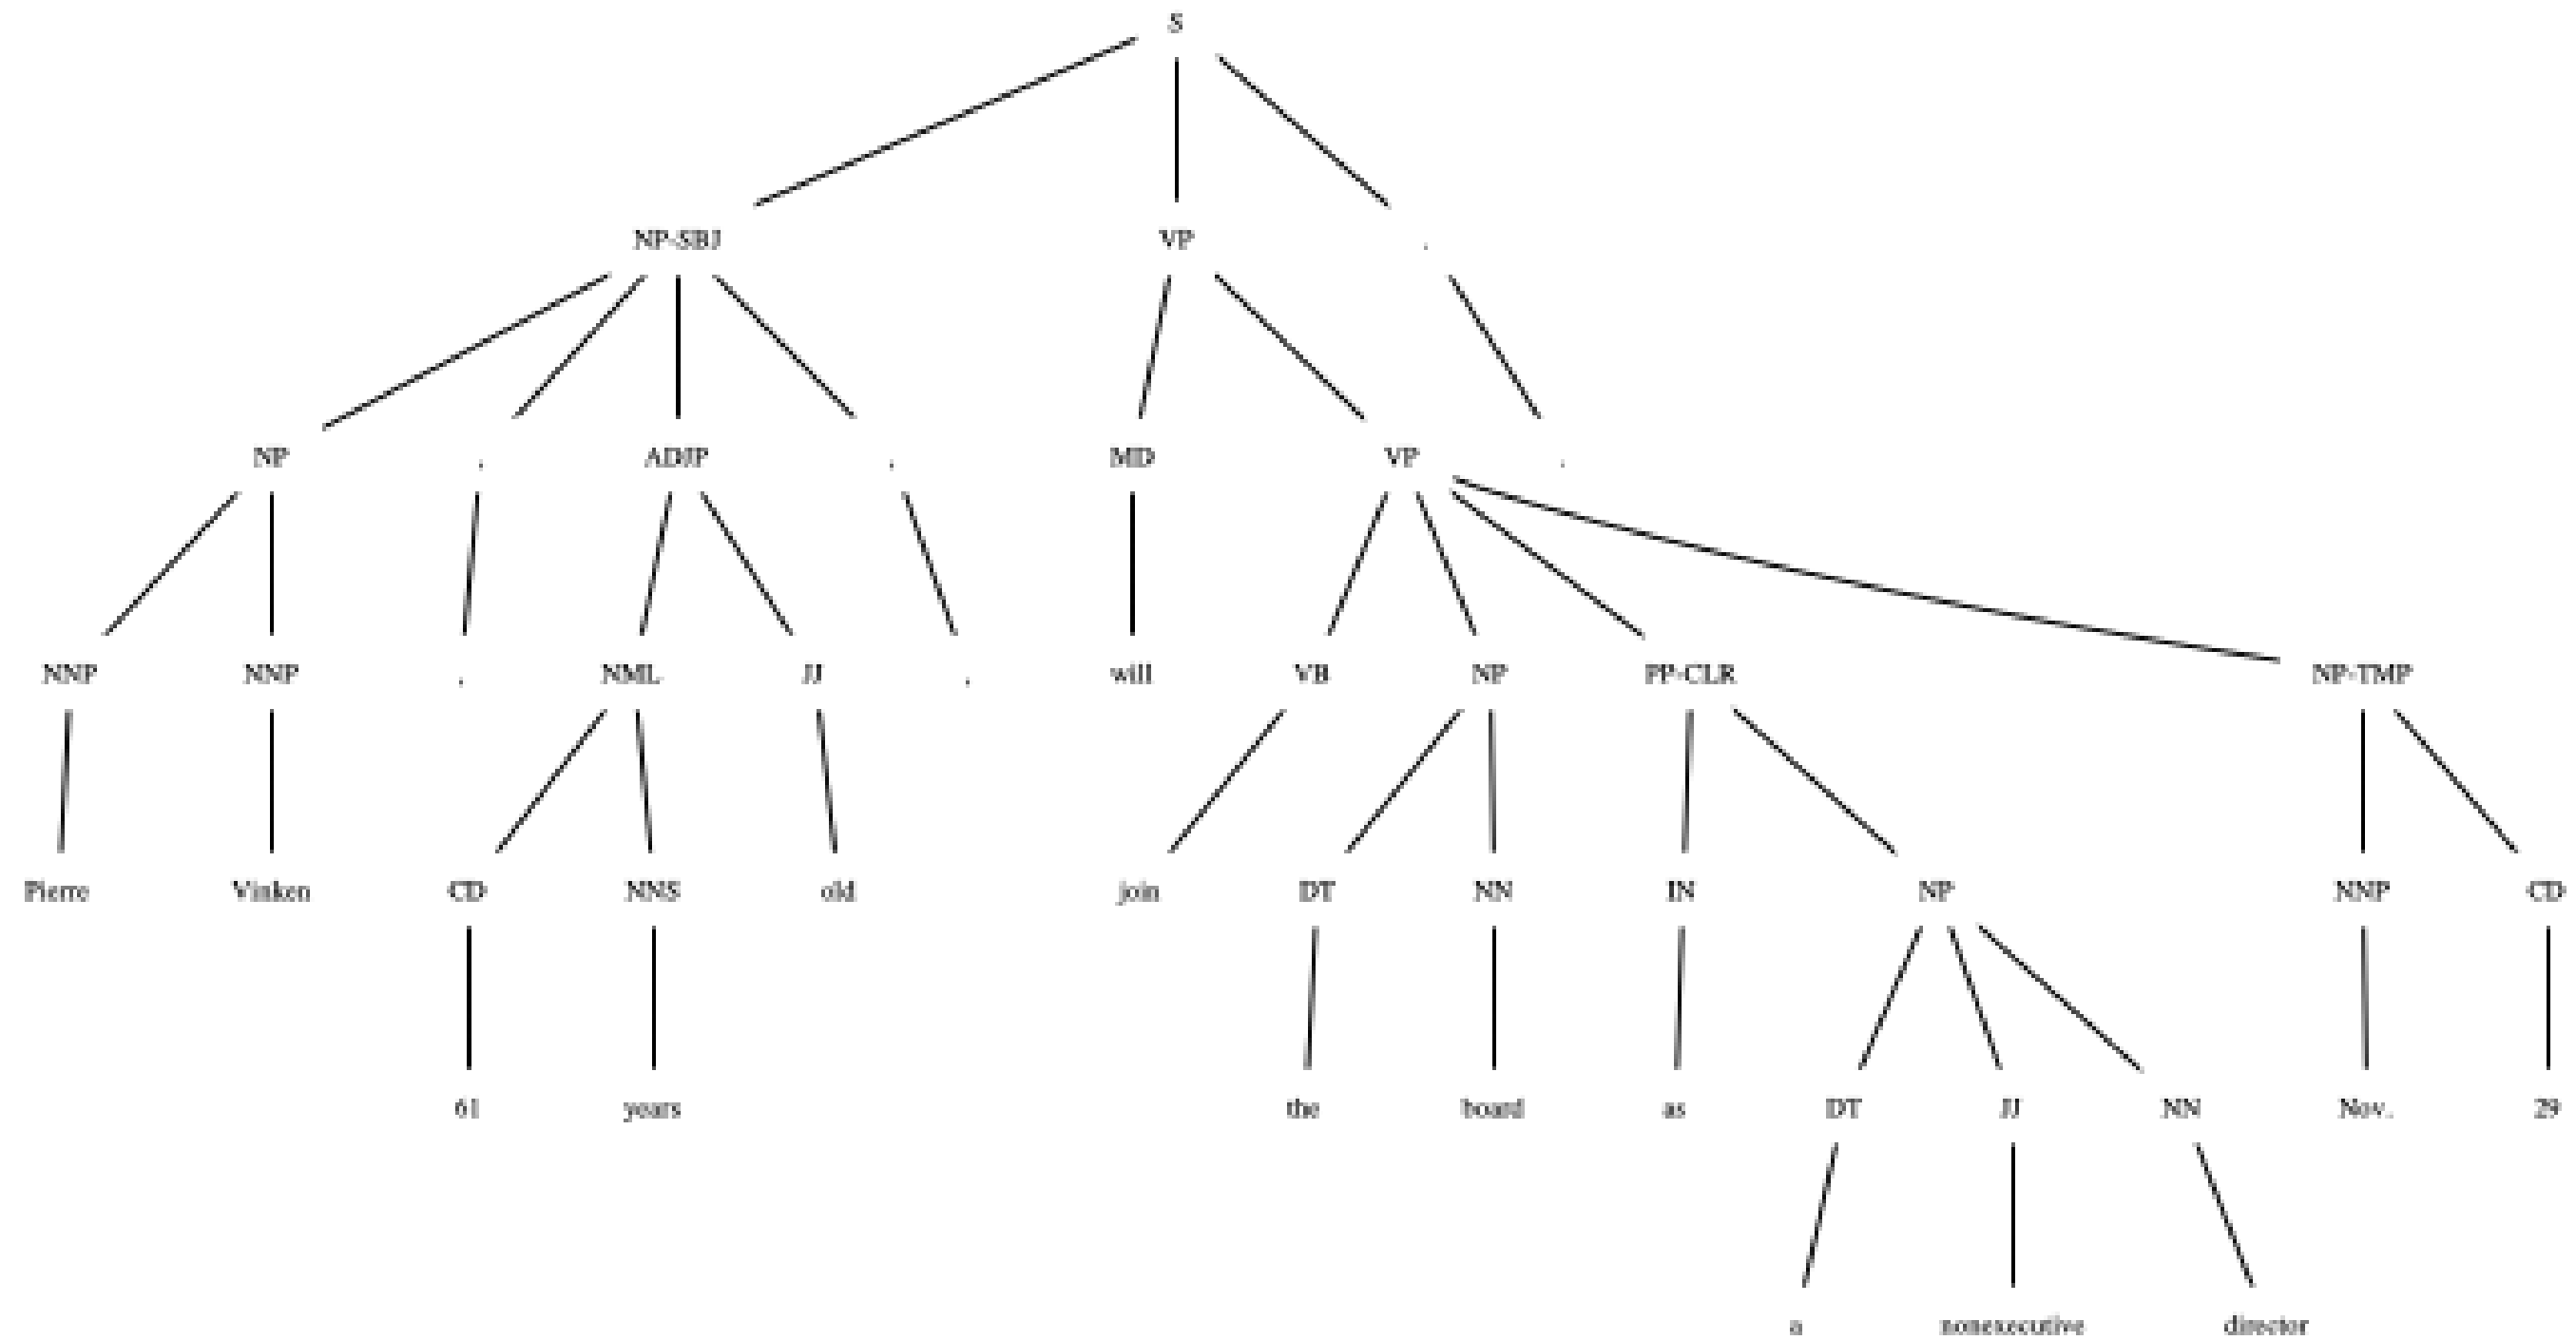
The Penn Treebank

- Syntactically annotated corpus of newspaper text (1989 Wall Street Journal Articles).
- The source text is naturally occurring but the treebank is not:
 - Assumes a specific linguistic theory (although a simple one).
 - Very flat structure (NPs, Ss, VPs).

PTB Example

```
( (S (NP-SBJ (NP (NNP Pierre) (NNP Vinken))
      (, ,)
      (ADJP (NML (CD 61) (NNS years))
      (JJ old))
      (, ,))
  (VP (MD will)
    (VP (VB join)
      (NP (DT the) (NN board))
      (PP-CLR (IN as)
        (NP (DT a) (JJ nonexecutive) (NN director)))
      (NP-TMP (NNP Nov.) (CD 29))))
  (. .)))
```

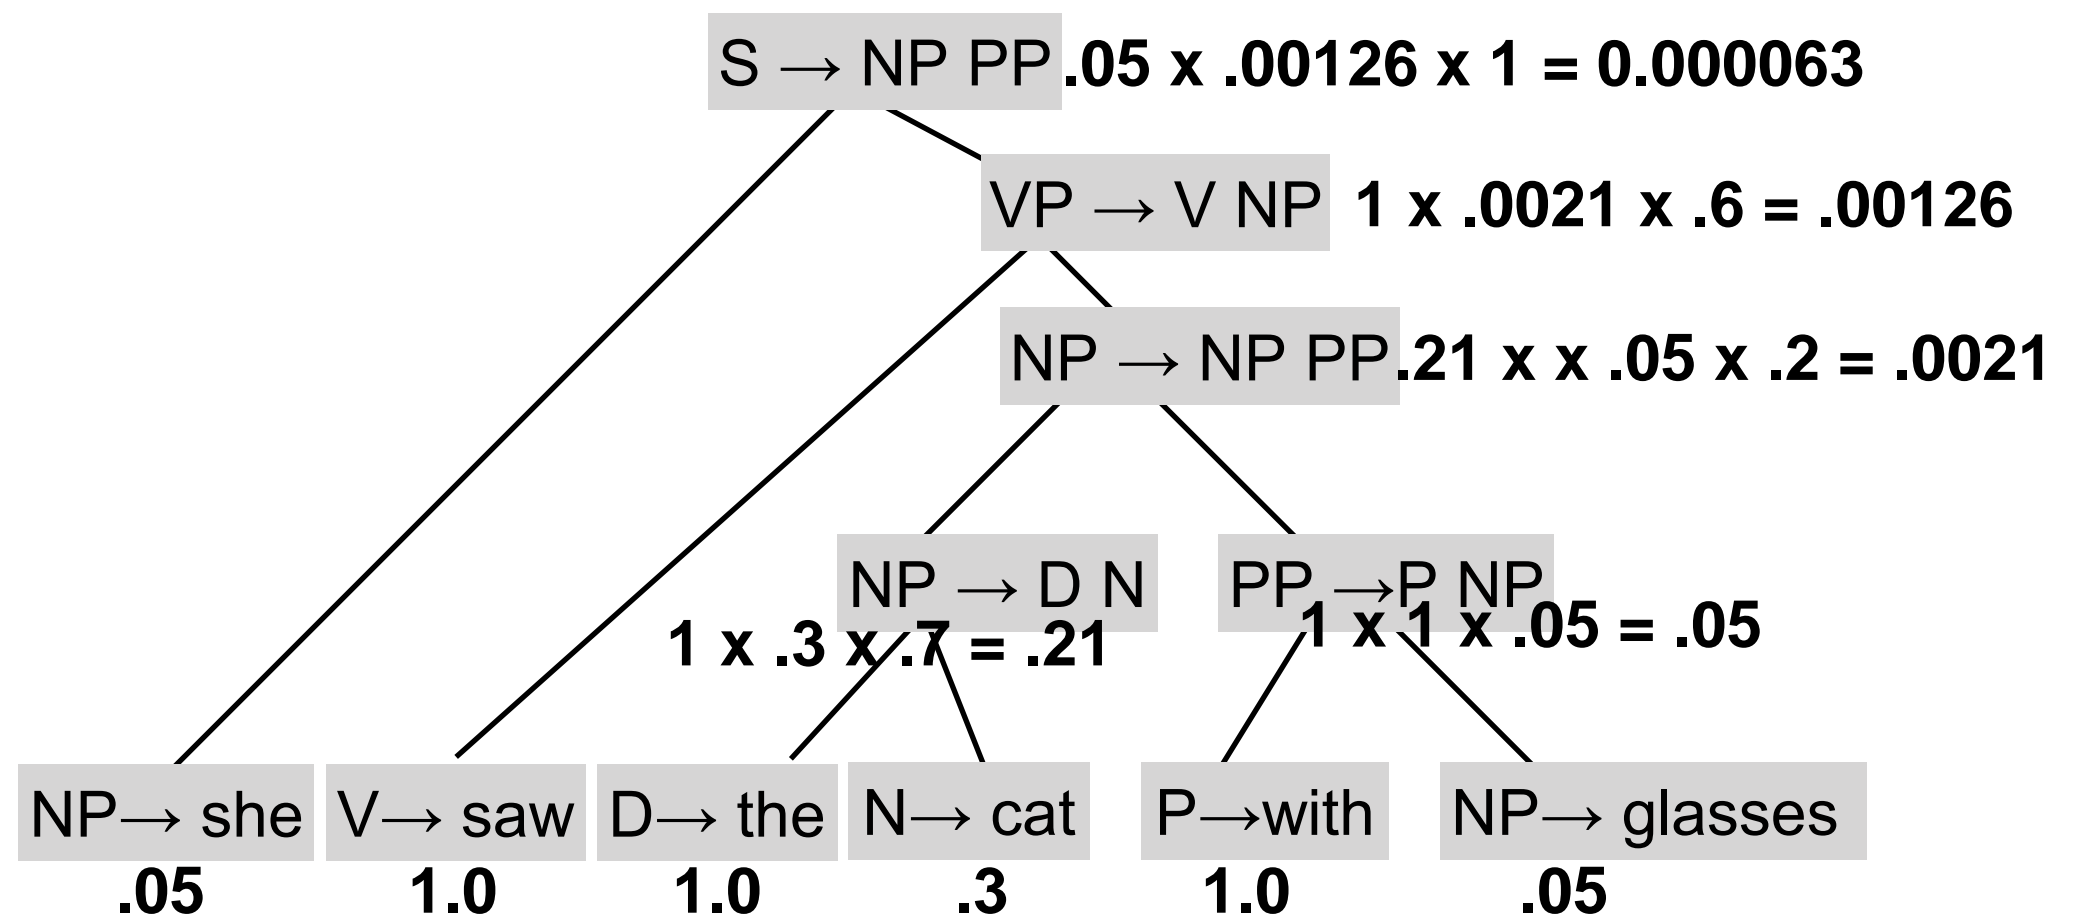
PTB Example



Parsing with PCFG

- We want to use PCFG to answer the following questions:
 - What is the total probability of the sentence under the PCFG?
 - What is the most probable parse tree for a sentence under the PCFG? (*decoding/parsing*)
- We can modify the CKY algorithm.
Basic idea: Compute these probabilities bottom-up using dynamic programming.

Computing Probabilities Bottom-Up



CKY for PCFG Parsing

- Let $T_G(s, A)$ be the set of trees generated by grammar G starting at nonterminal A , whose *yield* is string s
- Use a chart π so that $\pi[i, j, A]$ contains the probability of the highest probability parse tree for string $s[i, j]$ starting in nonterminal A .

$$\pi[i, j, A] = \max_{t \in T_G(s[i, j], A)} P(t)$$

- We want to find $\pi[0, \text{length}(s), S]$ -- the probability of the highest-scoring parse tree for s rooted in the start symbol S .

CKY for PCFG Parsing

- To compute $\pi[0, \text{length}(s), S]$ we can use the following recursive definition:

Base case:
$$\pi[i, i + 1, A] = \begin{cases} P(A \rightarrow s_i) & \text{if } A \rightarrow s_i \in R \\ 0 & \text{otherwise} \end{cases}$$

$$\pi[i, j, A] = \max_{\substack{k=i+1 \dots j-1, \\ A \rightarrow BC \in R}} P(A \rightarrow BC) \cdot \pi[i, k, B] \cdot \pi[k, j, C]$$

- Then fill the chart using dynamic programming.

CKY for PCFG Parsing

- **Input:** PCFG $G=(N, \Sigma, R, S)$, input string s of length n .
- for $i=0\dots n-1$: initialization
$$\pi[i, i+1, A] = \begin{cases} P(A \rightarrow s_i) & \text{if } A \rightarrow s_i \in R \\ 0 & \text{otherwise} \end{cases}$$
- for $length=2\dots n$: main loop
 - for $i=0\dots(n-length)$:
 - $j = i+length$
 - for $k=i+1\dots j-1$:
 - for $A \in N$:
 - $$\pi[i, j, A] = \max_{\substack{k=i+1\dots j-1, \\ A \rightarrow BC \in R}} P(A \rightarrow BC) \cdot \pi[i, k, B] \cdot \pi[k, j, C]$$

Use **backpointers** to retrieve the highest-scoring parse tree (see previous lecture).

Probability of a Sentence

- What if we are interested in the probability of a sentence, **not** of a single parse tree (for example, because we want to use the PCFG as a language model).
- Problem: Spurious ambiguity. Need to sum the probabilities of **all** parse trees for the sentence.
- How do we have to change CKY to compute this?

$$\pi[i, j, A] = \sum_{\substack{k=i+1 \dots j-1, \\ A \rightarrow BC \in R}} P(A \rightarrow BC) \cdot \pi[i, k, B] \cdot \pi[k, j, C]$$

Earley Parser

- CKY parser starts with words and builds parse trees bottom-up; requires the grammar to be in CNF.
- The Earley parser instead starts at the start symbol and tries to “guess” derivations top-down.
 - It discards derivations that are incompatible with the sentence.
 - The early parser sweeps through the sentence left-to-right only once. It keeps partial derivations in a table (“chart”).
 - Allows arbitrary CFGs, no limitation to CNF.

Parser States

- Earley parser keeps track of partial derivations using parser **states / items**.
- State represent hypotheses about constituent structure based on the grammar, taking into account the input.
- Parser states are represented as **dotted rules with spans**.
 - The constituents to the left of the \cdot have already been seen in the input string s (corresponding to the span)

$S \rightarrow \cdot NP VP [0,0]$ *“According to the grammar, there may be an NP starting in position 0. “*

$NP \rightarrow D A \cdot N [0,2]$ *“There is a determiner followed by an adjective in $s[0,2]$ “*

$NP \rightarrow NP PP \cdot [3,8]$ *“There is a complete NP in $s[3,8]$, consisting of an NP and PP”*

Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

$S \rightarrow \cdot NP VP [0,0]$

$NP \rightarrow \cdot NP PP [0,0]$ $NP \rightarrow \cdot D N [0,0]$

$D \rightarrow \cdot \text{the} [0,0]$

Three parser operations:

1. **Predict** new subtrees top-down.

$_0$ *the* $_1$ *student* $_2$ *saw* $_3$ *the* $_4$ *cat* $_5$ *with* $_6$ *the* $_7$ *tail*

Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

$S \rightarrow \cdot NP VP [0,0]$

$NP \rightarrow \cdot NP PP [0,0]$

$NP \rightarrow \cdot D N [0,0]$

$D \rightarrow \text{the} \cdot [0,1]$

Three parser operations:

1. Predict new subtrees top-down.
2. **Scan** input terminals.

0 *the* 1 *student* 2 *saw* 3 *the* 4 *cat* 5 *with* 6 *the* 7 *tail*

Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

$S \rightarrow \cdot NP VP [0,0]$

$NP \rightarrow \cdot NP PP [0,0]$

$NP \rightarrow \cdot D N [0,0]$

$D \rightarrow \text{the} \cdot [0,1]$

passive state

Three parser operations:

1. Predict new subtrees top-down.
2. **Scan** input terminals.

0 *the* 1 *student* 2 *saw* 3 *the* 4 *cat* 5 *with* 6 *the* 7 *tail*

Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

$NP \rightarrow \cdot NP PP [0,0]$

$S \rightarrow \cdot NP VP [0,0]$

$NP \rightarrow D \cdot N [0,1]$

$D \rightarrow \text{the} \cdot [0,1]$

passive state

Three parser operations:

1. Predict new subtrees top-down.
2. Scan input terminals.
3. **Complete** with passive states.

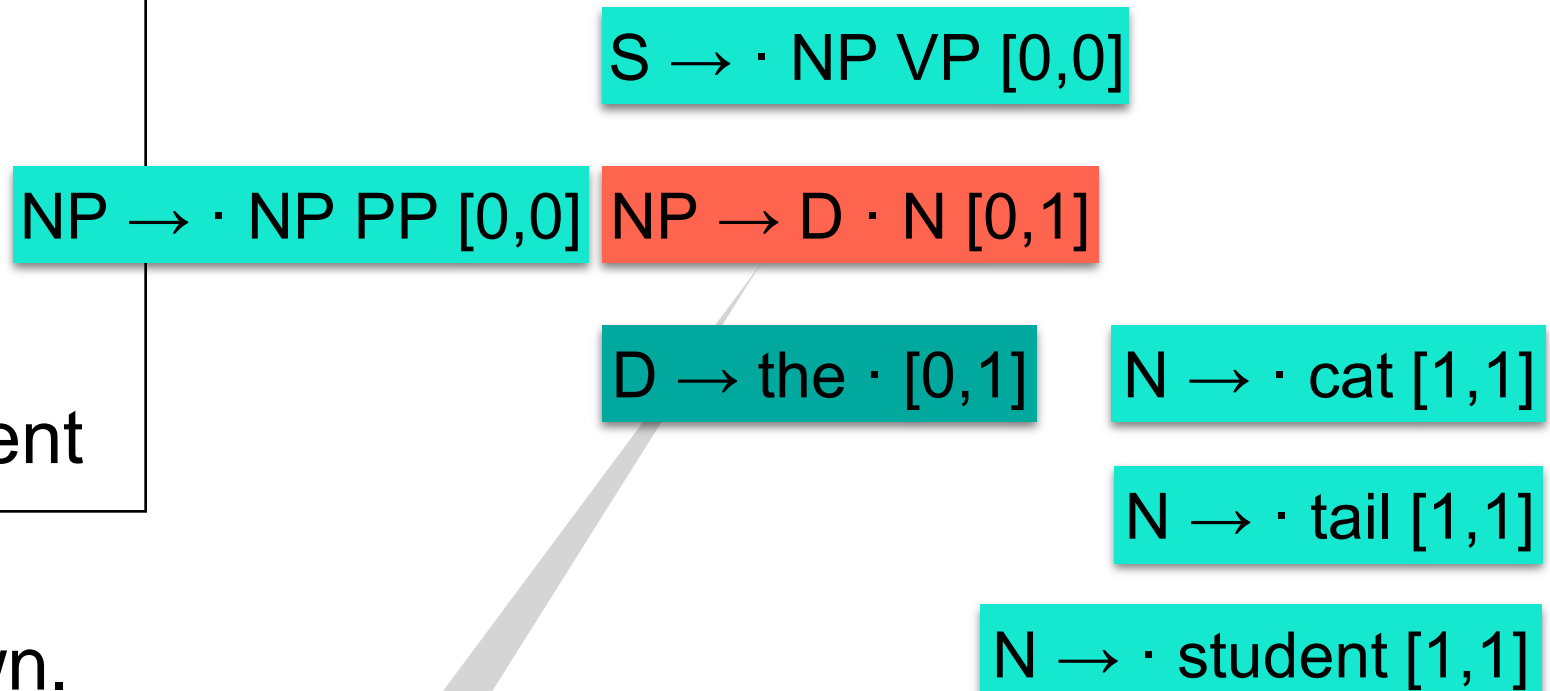
0 *the* 1 *student* 2 *saw* 3 *the* 4 *cat* 5 *with* 6 *the* 7 *tail*

Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

Three parser operations:

1. **Predict** new subtrees top-down.
2. Scan input terminals.
3. Complete with passive states.



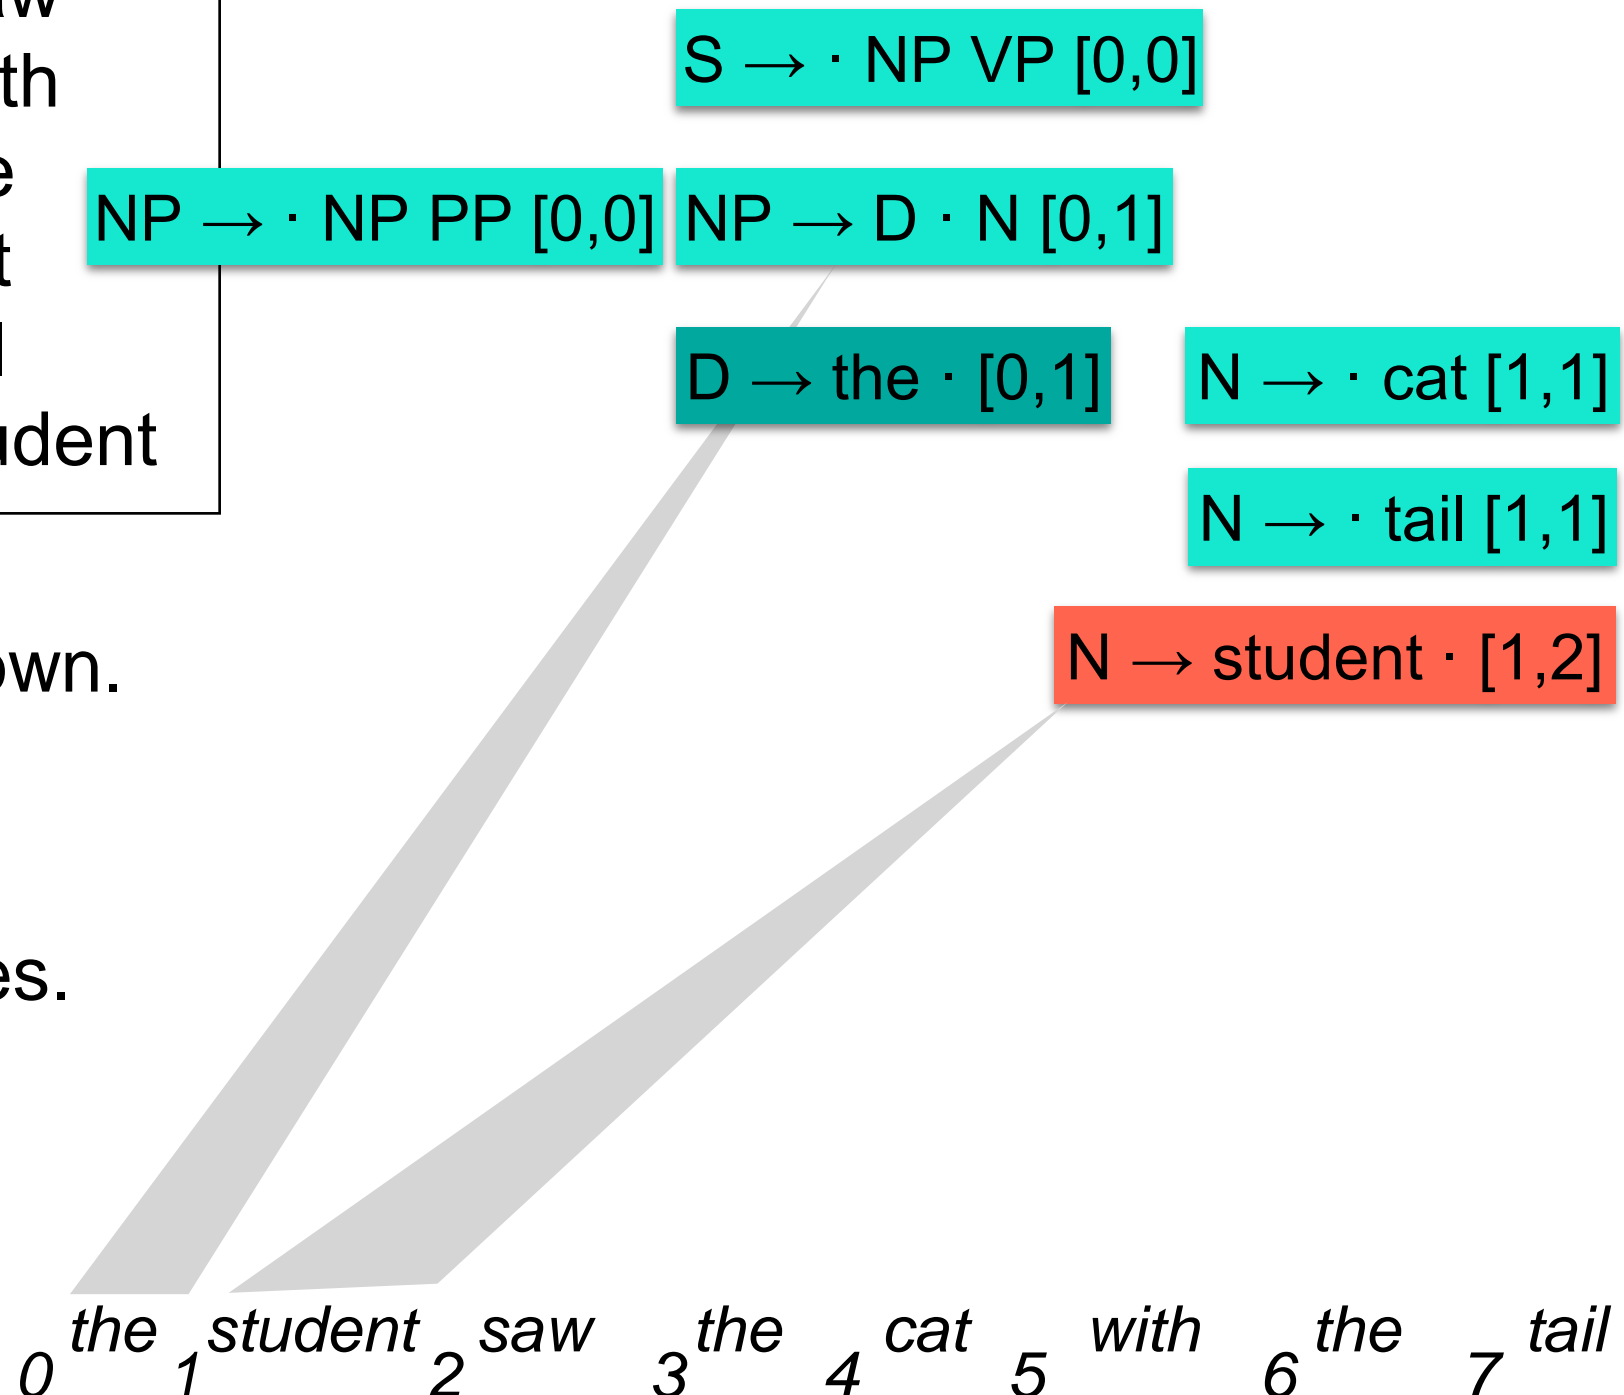
0 the 1 student 2 saw 3 the 4 cat 5 with 6 the 7 tail

Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

Three parser operations:

1. Predict new subtrees top-down.
2. **Scan** input terminals.
3. Complete with passive states.

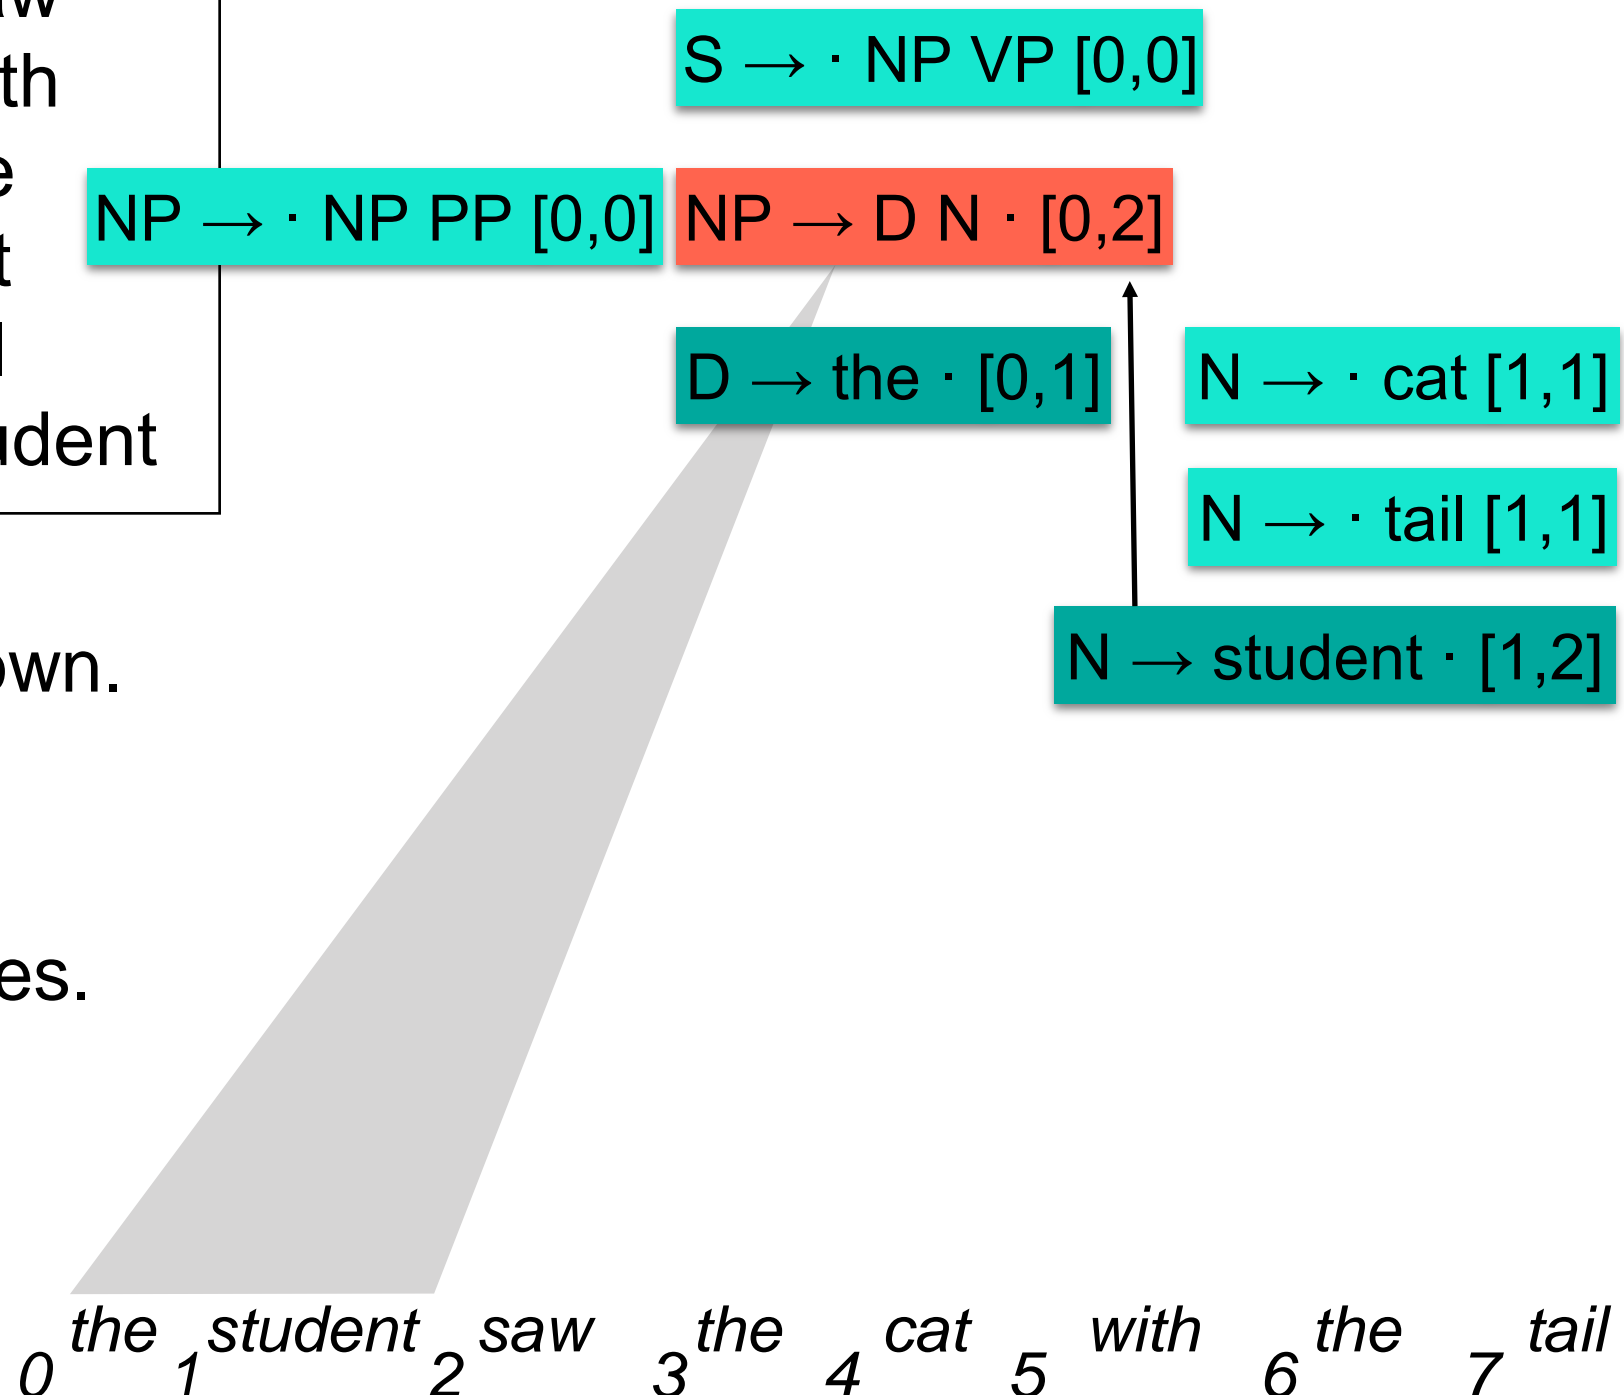


Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

Three parser operations:

1. Predict new subtrees top-down.
2. Scan input terminals.
3. **Complete** with passive states.

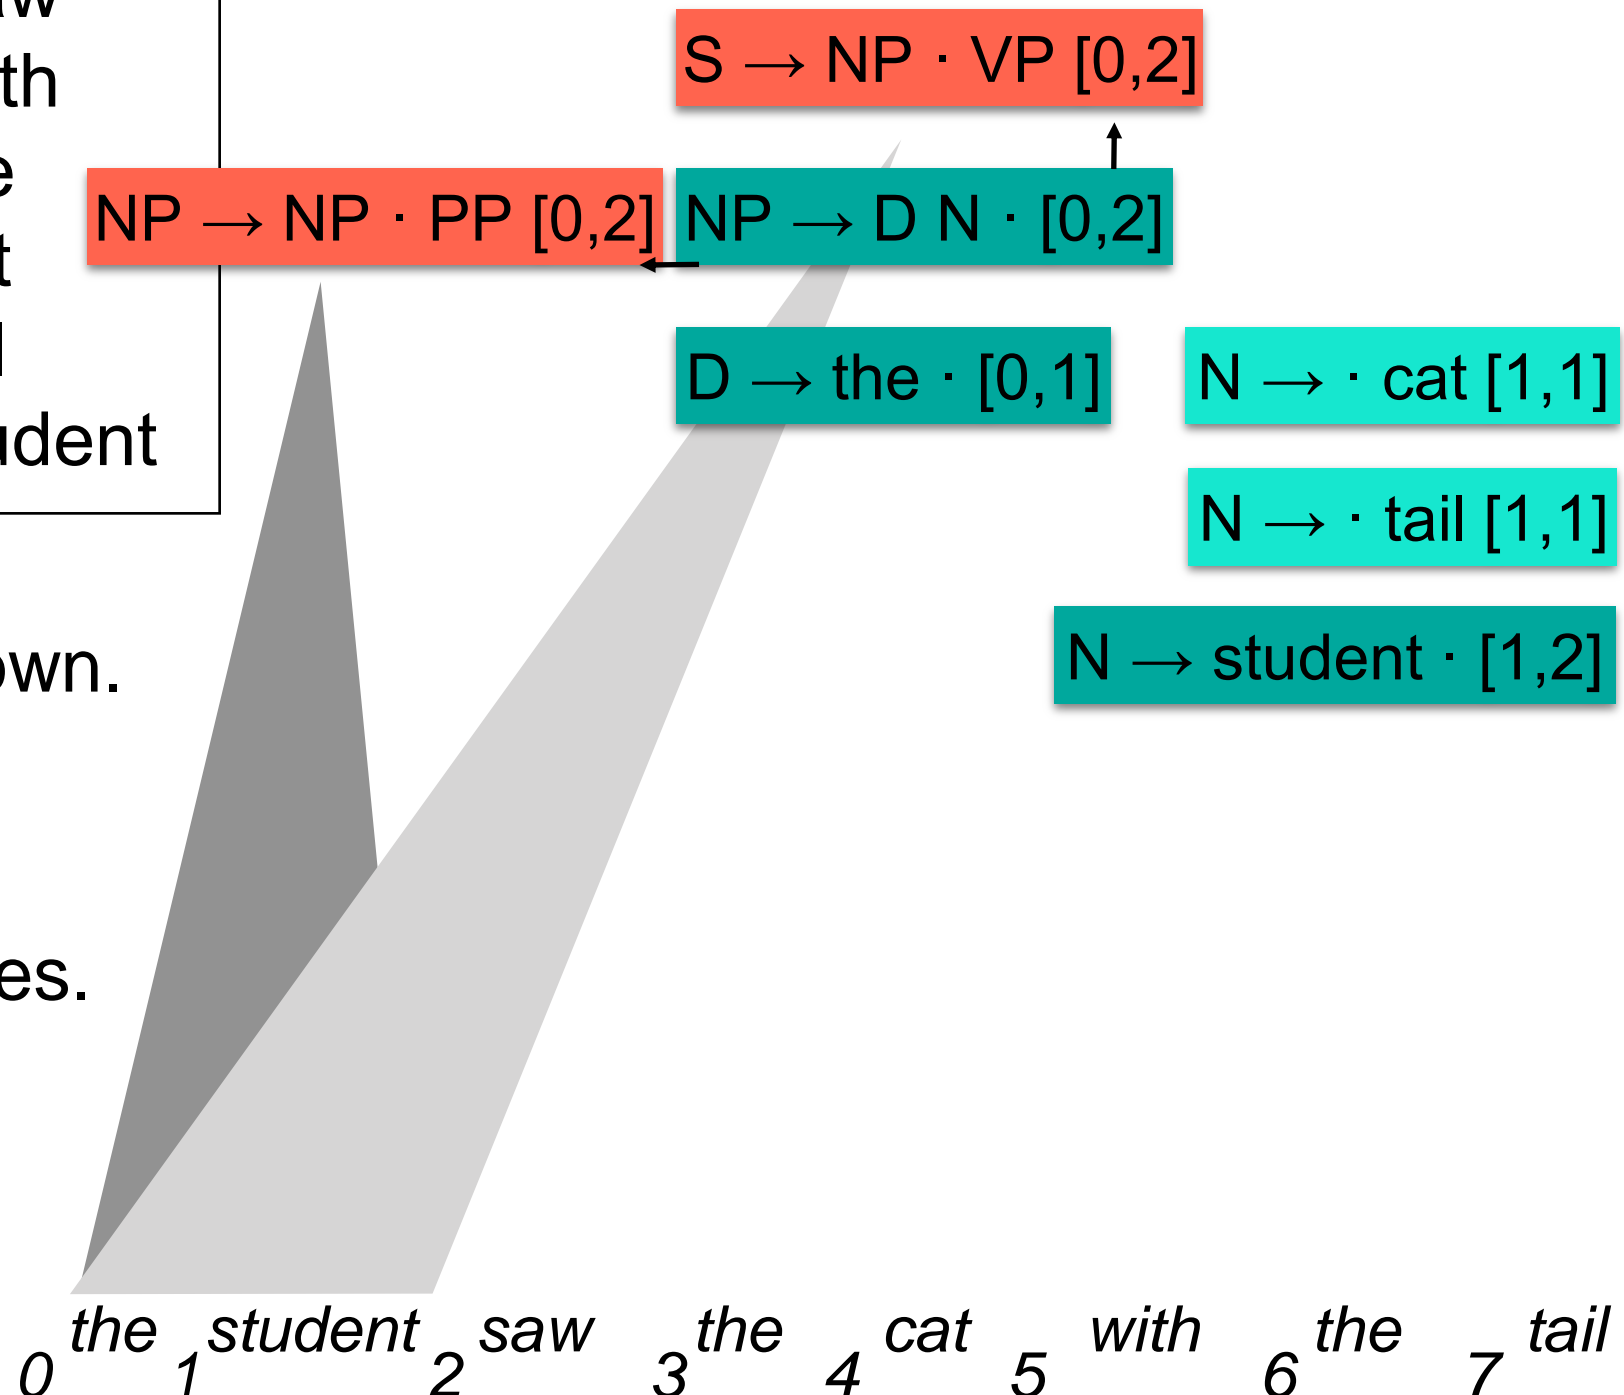


Earley Parser (sketch)

$S \rightarrow NP VP$	$V \rightarrow \text{saw}$
$VP \rightarrow V NP$	$P \rightarrow \text{with}$
$VP \rightarrow VP PP$	$D \rightarrow \text{the}$
$PP \rightarrow P NP$	$N \rightarrow \text{cat}$
$NP \rightarrow D N$	$N \rightarrow \text{tail}$
$NP \rightarrow NP PP$	$N \rightarrow \text{student}$

Three parser operations:

1. Predict new subtrees top-down.
2. Scan input terminals.
3. **Complete** with passive states.



Earley Algorithm

- Keep track of parser states in a table (“chart”). $Chart[k]$ contains a set of all parser states that end in position k .
- **Input:** Grammar $G=(N, \Sigma, R, S)$, input string s of length n .
- **Initialization:** For each production $S \rightarrow \alpha \in R$
add a state $S \rightarrow \cdot \alpha [0, 0]$ to $Chart[0]$.
- for $i = 0$ to n :
 - for each $state$ in $Chart[i]$:
 - if $state$ is of form $A \rightarrow \alpha \cdot s[i] \beta [k, i]$:
scan($state$)
 - elif $state$ is of form $A \rightarrow \alpha \cdot B \beta [k, i]$:
predict($state$)
 - elif $state$ is of form $A \rightarrow \alpha \cdot [k, i]$:
complete($state$)

Earley Algorithm

- Keep track of parser states in a table (“chart”). $Chart[k]$ contains a set of all parser states that end in position k .

- **Input:** Grammar $G=(N, \Sigma, R, S)$, input string s of length n .

- **Initialization:** For each production $S \rightarrow \alpha \in R$
add a state $S \rightarrow \cdot \alpha [0,0]$ to $Chart[0]$.

- for $i = 0$ to n :

- for each $state$ in $Chart[i]$:

- if $state$ is of form $A \rightarrow \alpha \cdot s[i] \beta [k,i]$:
scan($state$)

- elif $state$ is of form $A \rightarrow \alpha \cdot B \beta [k,i]$:
predict($state$)

- elif $state$ is of form $A \rightarrow \alpha \cdot [k,i]$
complete($state$)

else then is states of form
 $A \rightarrow \alpha \cdot \beta [k,i]$, i.e.

β is not $s[i]$, in which case we
don't want to do anything

Earley Algorithm - Scan

- The scan operation can only be applied to a state if the dot is in front of a terminal symbol that matches the next input terminal.

- function **scan**(state): *// state is of form $A \rightarrow \alpha \cdot s[i] \beta [k, i]$*
 - Add a new state $A \rightarrow \alpha s[i] \cdot \beta [k, i+1]$
to *Chart*[*i+1*]

Earley Algorithm - Predict

- The predict operation can only be applied to a state if the dot is in front of a non-terminal symbol.
- function **predict**(state): *// state is of form $A \rightarrow \alpha \cdot B \beta [k, i]$:*
 - Add a new state $B \rightarrow \cdot \gamma [i, i]$
to *Chart[i]*
- Note that this modifies Chart[i] **while** the algorithm is looping through it.
- No duplicate states are added (Chart[i] is a set)

Earley Algorithm - Complete

- The complete operation may only be applied to a passive item.

- function **complete**(state): *// state is of form $A \rightarrow \alpha \cdot [k,j]$*

- for each state $B \rightarrow \beta \cdot A \gamma [i,k]$ add a new state $B \rightarrow \beta A \cdot \gamma [i,j]$ to Chart[j]

- Note that this modifies Chart[i] **while** the algorithm is looping through it.
- Note that it is important to make a copy of the old state before moving the dot.
- This operation is similar to the combination operation in CKY!

Earley Algorithm - Runtime

- The runtime depends on the number of items in the chart (each item is “visited” exactly once).
- We proceed through the input exactly once, which takes $O(N)$.
- For each position on the chart, there are $O(N)$ possible split points where the dot could be.
- Each complete operation can produce $O(N)$ possible new items (with different starting points).
- Total: $O(N^3)$

Earley Algorithm - Some Observations

- How do we recover parse trees?
 - What happens in case of ambiguity?
 - Multiple ways to Complete the same state.
 - Keep back-pointers in the parser state objects.
 - Or use a separate data structure (CKY-style table or hashed states)
- How do we make the algorithm work with PCFG?
 - Easy to compute probabilities on Complete. Follow back pointer with max probability.